



Processes and Threads

경희대학교 컴퓨터공학과

조진성

Program vs. Process

■ Program

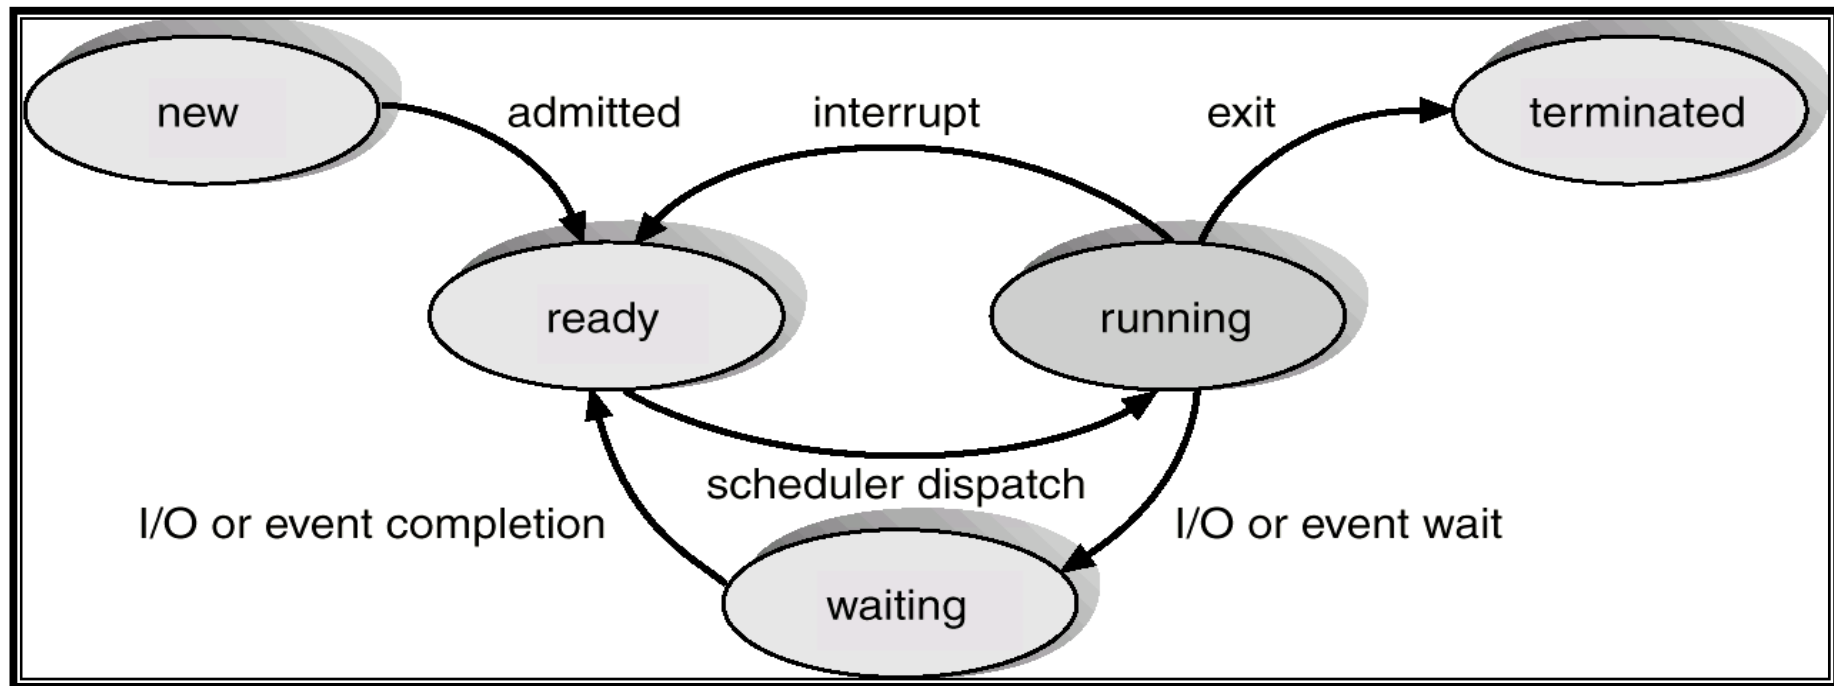
- ✓ Executable file on a disk
- ✓ Loaded into memory and executed by the kernel

■ Process

- ✓ Executing instance of a program
- ✓ The basic unit of execution and scheduling
- ✓ A process is named using its process ID (PID)
- ✓ Other IDs associated with a process
 - real user ID
 - real group ID
 - effective user ID
 - effective group ID
 - saved set-user-ID
 - saved set-group-ID

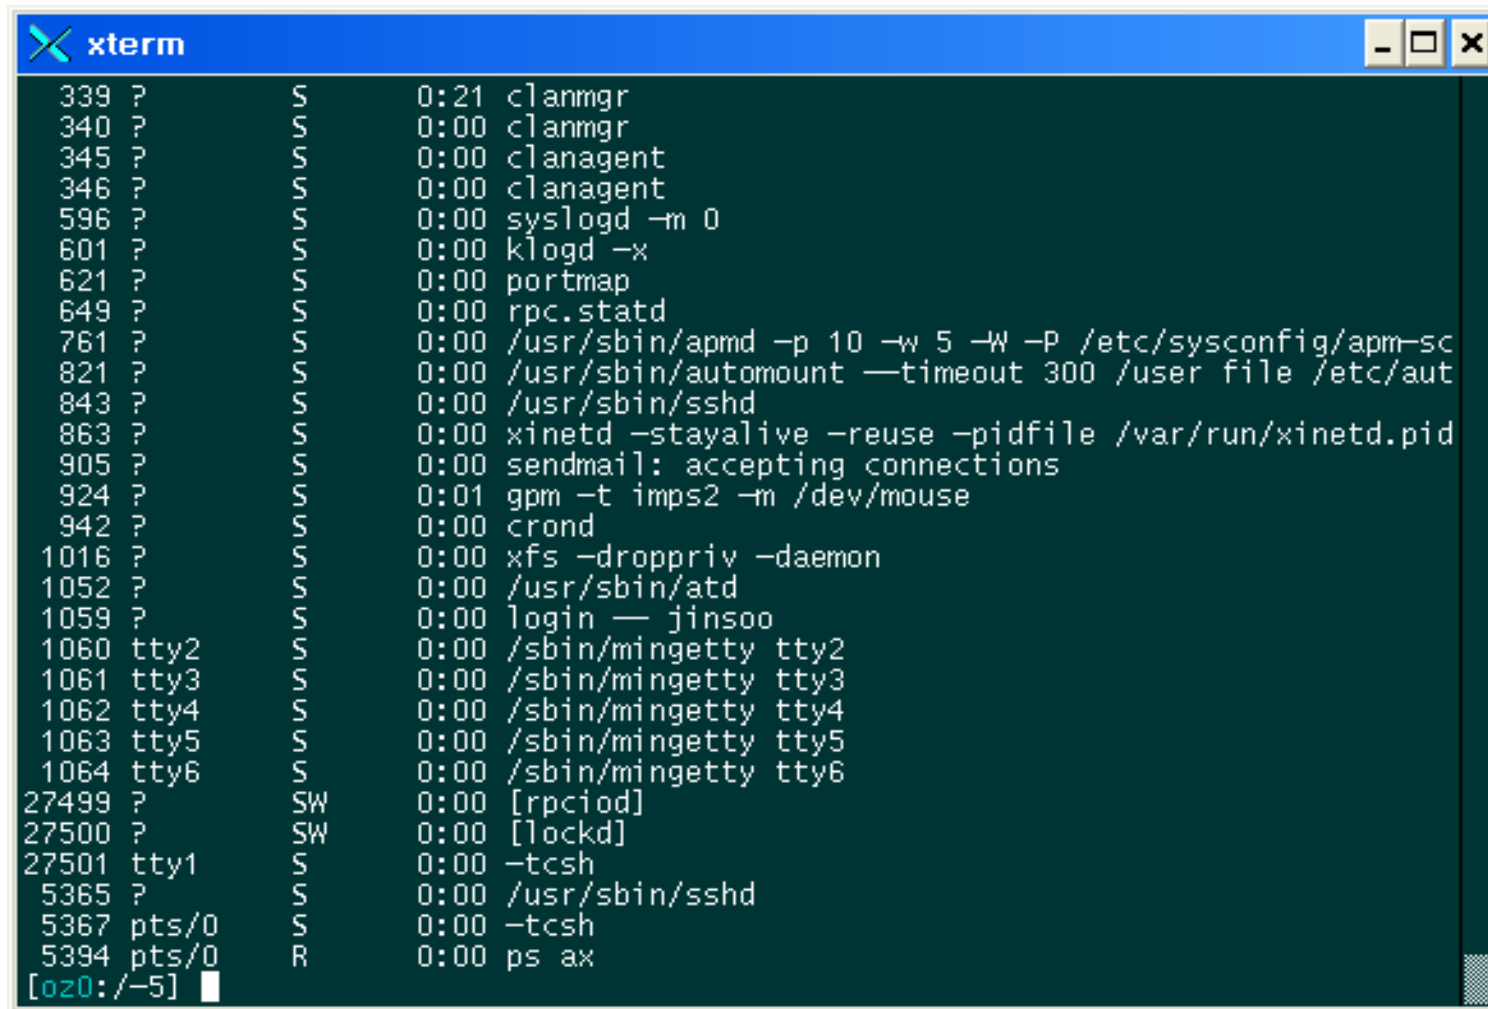


Process State



Process State (Cont'd)

■ “ps” command



```
xterm
339 ?      S      0:21  clanmgr
340 ?      S      0:00  clanmgr
345 ?      S      0:00  clanagent
346 ?      S      0:00  clanagent
596 ?      S      0:00  syslogd -m 0
601 ?      S      0:00  klogd -x
621 ?      S      0:00  portmap
649 ?      S      0:00  rpc.statd
761 ?      S      0:00  /usr/sbin/apmd -p 10 -w 5 -W -P /etc/sysconfig/apm-sc
821 ?      S      0:00  /usr/sbin/automount --timeout 300 /user file /etc/aut
843 ?      S      0:00  /usr/sbin/sshd
863 ?      S      0:00  xinetd -stayalive -reuse -pidfile /var/run/xinetd.pid
905 ?      S      0:00  sendmail: accepting connections
924 ?      S      0:01  gpm -t imps2 -m /dev/mouse
942 ?      S      0:00  crond
1016 ?     S      0:00  xfs -droppriv -daemon
1052 ?     S      0:00  /usr/sbin/atd
1059 ?     S      0:00  login -- jinsoo
1060 tty2   S      0:00  /sbin/minigetty tty2
1061 tty3   S      0:00  /sbin/minigetty tty3
1062 tty4   S      0:00  /sbin/minigetty tty4
1063 tty5   S      0:00  /sbin/minigetty tty5
1064 tty6   S      0:00  /sbin/minigetty tty6
27499 ?    SW     0:00  [rpciod]
27500 ?    SW     0:00  [lockd]
27501 tty1  S      0:00  -tcsh
5365 ?     S      0:00  /usr/sbin/sshd
5367 pts/0  S      0:00  -tcsh
5394 pts/0  R      0:00  ps ax
[oz0:/-5]
```

- R: Runnable
- S: Sleeping
- T: Traced or Stopped
- D: Uninterruptible Sleep
- Z: Zombie

- W: No resident pages
- <: High-priority task
- N: Low-priority task
- L: Has pages locked into memory



IDs Associated with a Process

■ Get various IDs associated with a process

- ✓ `#include <sys/types.h>`
- ✓ `#include <unistd.h>`
- ✓ `pid_t getpid(void);`
- ✓ return: process ID of calling process
- ✓ `pid_t getppid(void);`
- ✓ return: parent process ID of calling process
- ✓ `uid_t getuid(void);`
- ✓ return: real user ID of calling process
- ✓ `uid_t geteuid(void);`
- ✓ return: effective user ID of calling process
- ✓ `gid_t getgid(void);`
- ✓ return: group ID of calling process
- ✓ `gid_t getegid(void);`
- ✓ return: effective group ID of calling process



Create a New Process

■ fork: the only way a new process is created in Linux

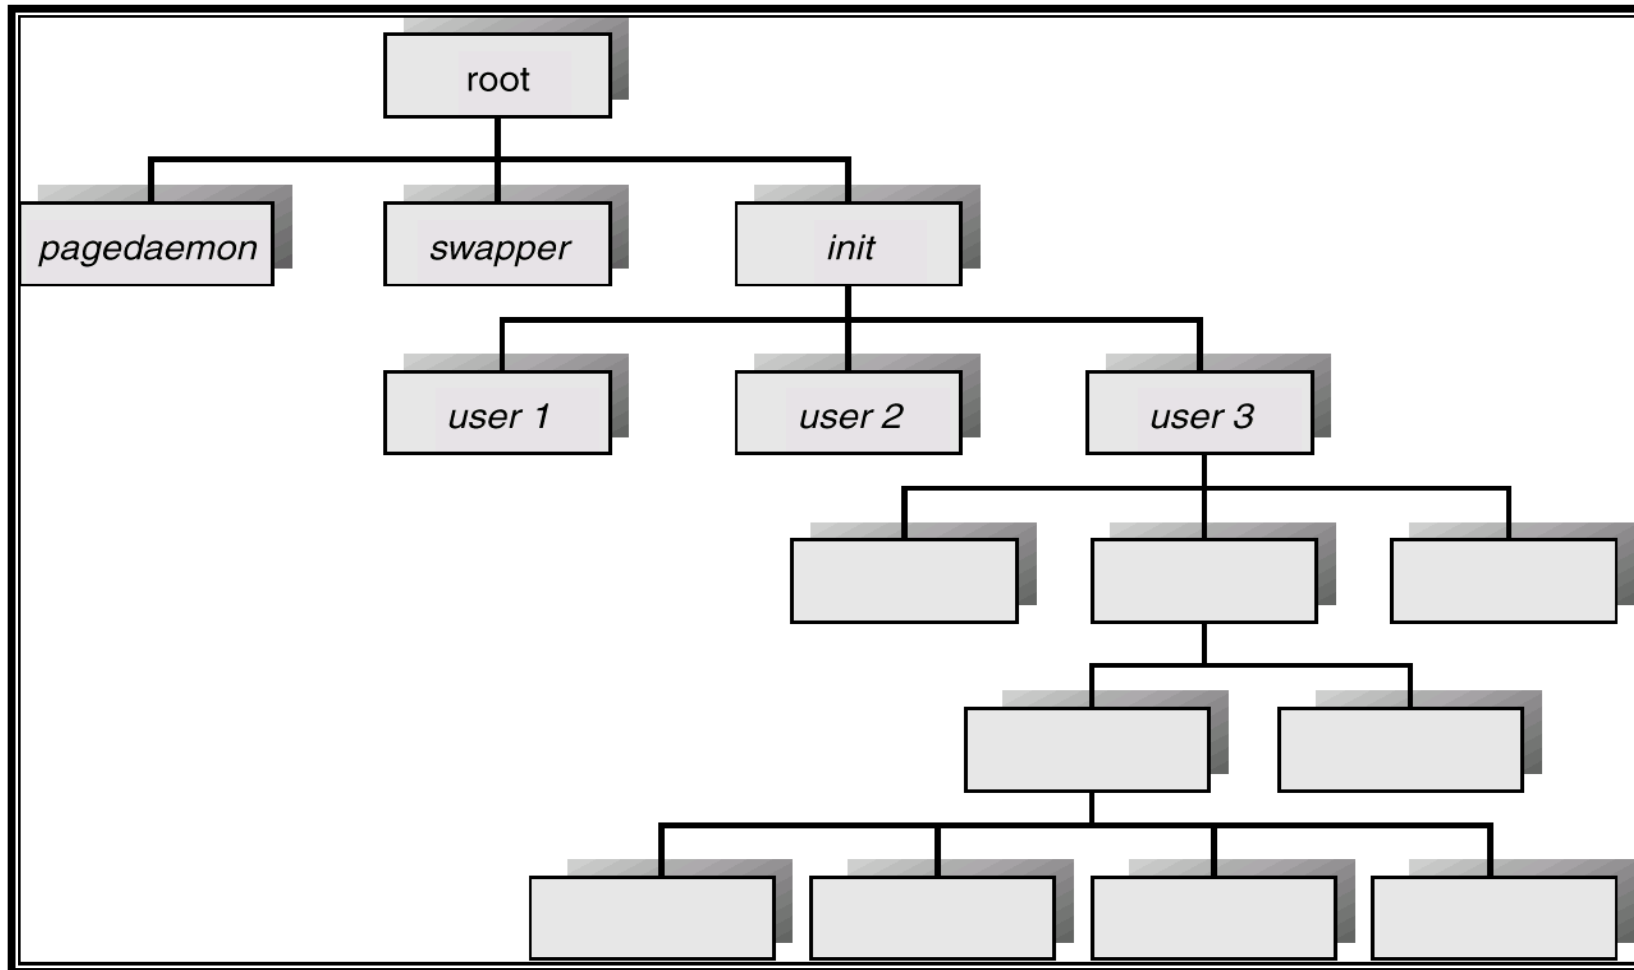
- ✓ `#include <sys/types.h>`
- ✓ `#include <unistd.h>`
- ✓ `pid_t fork(void);`
- ✓ return: 0 in child, process ID of child in parent, -1 on error

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
main()
{
    int pid;
    if ((pid = fork()) == 0)
        /* child */
        printf("I am %d. My parent is %d\n", getpid(), getppid());
    else
        /* parent */
        printf("I am %d. My child is %d\n", getpid(), pid);
}
```



Create a New Process (Cont'd)

■ Process tree in Linux



Create a New Process (Cont'd)

- Why fork() ? → Very useful when the child...
 - ✓ is cooperating with the parent
 - ✓ relies upon the parent's data to accomplish its task
 - ✓ Example: Web server

```
while (1) {  
    int sock = accept();  
    if ((pid = fork()) == 0) {  
        /* Handle client request */  
    } else {  
        /* Close socket */  
    }  
}
```



Create a New Process (Cont'd)

- Sharing of open files between parent and child after `fork`

Parent Process Table

	fd flags	ptr
fd 0 :		
fd 1 :		
fd 2 :		

Child Process Table

	fd flags	ptr
fd 0 :		
fd 1 :		
fd 2 :		

File Table

File status flags
Current file offset
v-node ptr

File status flags
Current file offset
v-node ptr

File status flags
Current file offset
v-node ptr

v-node Table

v-node information
i-node information

v-node information
i-node information

v-node information
i-node information



Exercise

■ fork example

```
$ gcc -o fork fork.c (or make fork)
```

```
$ ./fork
```



Terminate a Process

■ Normal termination

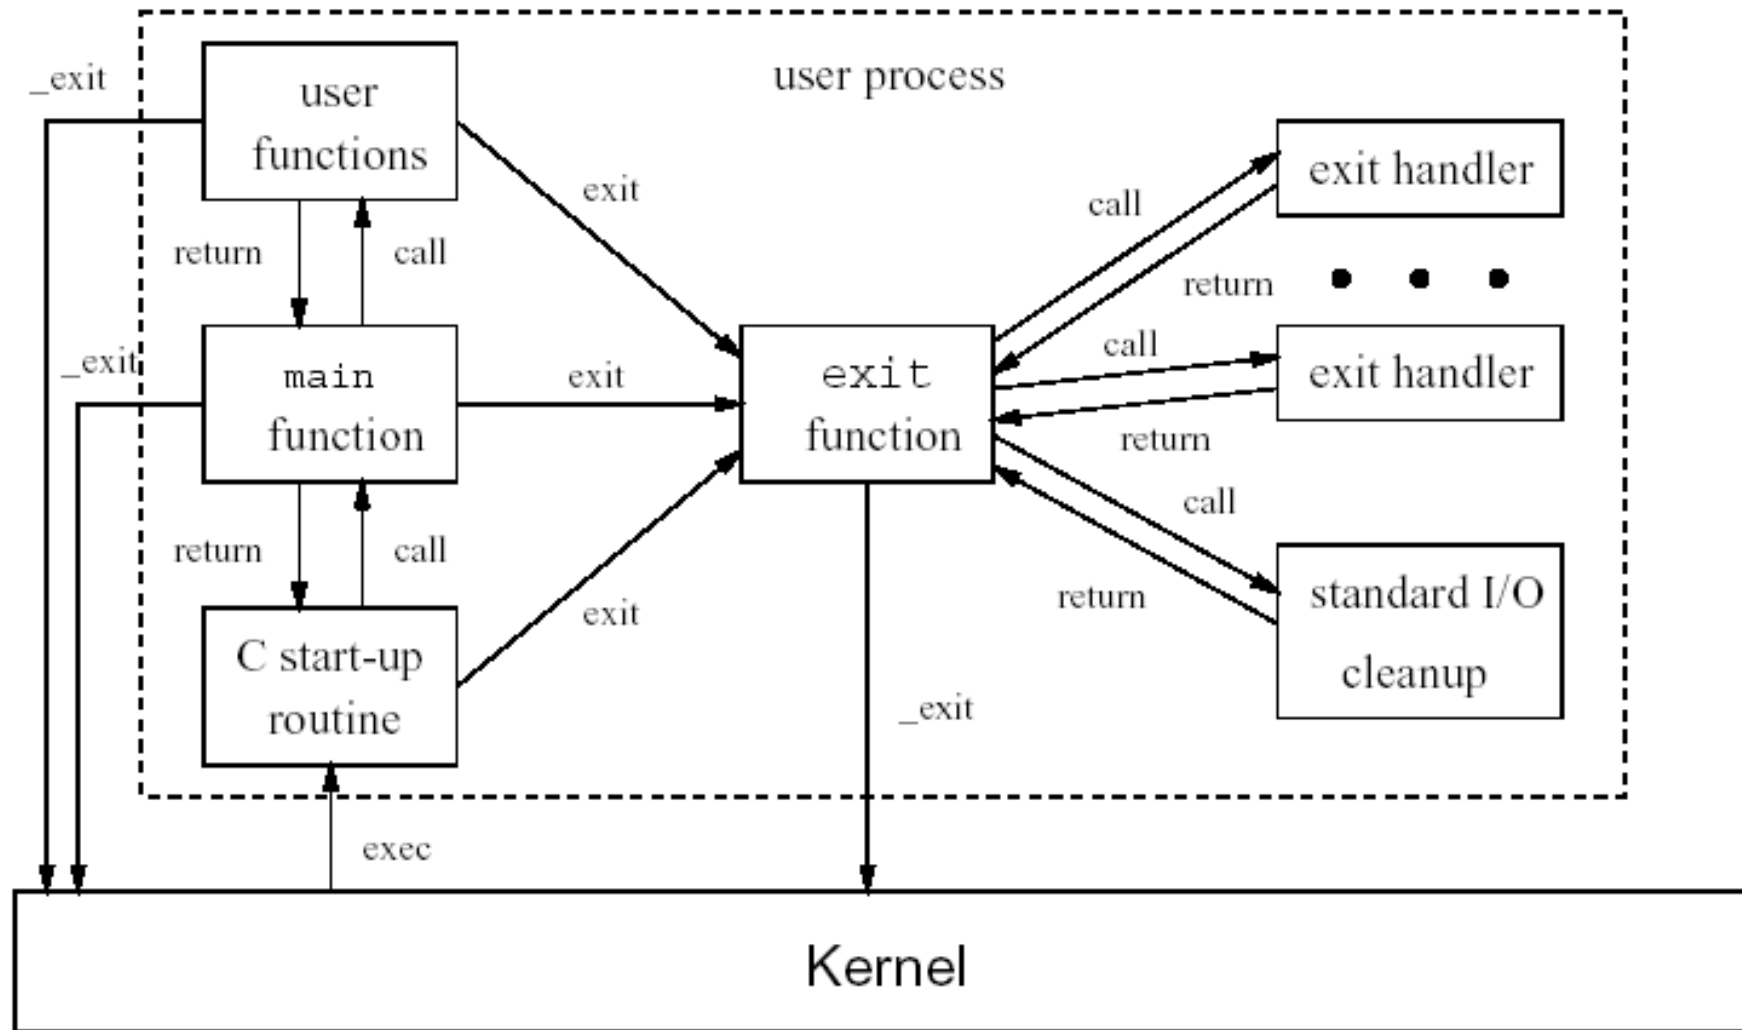
- ✓ return from `main()`
- ✓ calling `exit()`
- ✓ calling `_exit()`

■ Abnormal termination

- ✓ calling `abort()`
- ✓ terminated by a signal



Start and Termination of a C Program



Terminate a Process

■ exit

- ✓ `#include <stdlib.h>`
- ✓ `void exit(int status);`
- ✓ return: 0 if OK, nonzero on error

- ✓ `#include <unistd.h>`
- ✓ `void _exit(int status);`
- ✓ return: 0 if OK, nonzero on error

■ Register an exit handler

- ✓ `#include <stdlib.h>`
- ✓ `int atexit(void (*func)(void));`
- ✓ return: 0 if OK, nonzero on error



Exercise

■ atexit example

```
$ gcc -o exit exit.c (or make exit)
```

```
$ ./exit
```



Wait for Process Termination

■ wait

- ✓ `#include <sys/types.h>`
- ✓ `#include <sys/wait.h>`
- ✓ `pid_t wait(int *statloc);`
- ✓ `pid_t waitpid(pid_t pid, int *statloc, int options);`
- ✓ return: process ID if OK, 0 (see below) or -1 on error
 - With `WNOHANG` option, `waitpid` will not block if a child specified by `pid` is not immediately available. In this case, the return value is 0
- ✓ The calling process will
 - block (if all of its children are still running)
 - return immediately with the termination status of a child (if a child has terminated and is waiting for its termination status to be fetched)
 - return immediately with an error (if it doesn't have any child processes)

■ Cf) `SIGCHLD` signal

- ✓ asynchronous event



Exercise

■ wait example

```
$ gcc -o wait wait.c (or make wait)
```

```
$ ./wait
```

■ A program with race condition

```
$ gcc -o race race.c (or make race)
```

```
$ ./race
```

■ Modification to avoid race condition using `wait` system call

```
$ gcc -o worace worace.c (or make worace)
```

```
$ ./worace
```



Execute Another Program in a Program

■ exec

- ✓ `#include <unistd.h>`
- ✓ `int execl(char *pathname, char *arg0, ... /* (char *)0 */);`
- ✓ `int execv(char *pathname, char *argv[]);`
- ✓ `int execlp(char *pathname, char *arg0, ... /* (char *)0, char *envp[] */);`
- ✓ `int execve(char *pathname, char *argv[], char *envp[]);`
- ✓ `int execlp(char *filename, char *arg0, ... /* (char *)0 */);`
- ✓ `int execvp(char *filename, char *argv[]);`
- ✓ return: -1 on error, no return on success



Execute a Command String in a Program

■ system

- ✓ `#include <stdlib.h>`
- ✓ `int system(char *cmdstring);`
- ✓ return: termination status of the shell if OK, -1 on error
- ✓ `system` is implemented by calling `fork`, `exec`, and `waitpid`

```
#include <stdio.h>
#include <stdlib.h>
main()
{
    system("ls -al");
    system("date");
    system("who");
}
```



Exercise

■ Access environment variables

```
$ gcc -o env env.c (or make env)
```

```
$ ./env
```

■ exec example

```
$ gcc -o exec exec.c (or make exec)
```

```
$ ./exec
```

■ system example

```
$ gcc -o system system.c (or make system)
```

```
$ ./system
```



Process vs. Thread

■ Why thread ?

✓ Web server example using process

- Using fork() to create new processes to handle requests in parallel is overkill for such a simple task

```
while (1) {  
    int sock = accept();  
    if ((pid = fork()) == 0) {  
        /* Handle client request */  
    } else {  
        /* Close socket */  
    }  
}
```



Process vs. Thread (Cont'd)

■ Why thread ? (Cont'd)

- ✓ Web server example using thread
 - We can create a new thread for each request

```
webserver ()
{
    while (1) {
        int sock = accept();
        thread_fork(handle_request, sock);
    }
}

handle_request (int sock)
{
    /* Process request */
    close (sock);
}
```



Process vs. Thread (Cont'd)

■ Why thread ? (Cont'd)

- ✓ Responsiveness
- ✓ Resource Sharing
- ✓ Economy
- ✓ Utilization of MP Architectures





Thread Concept

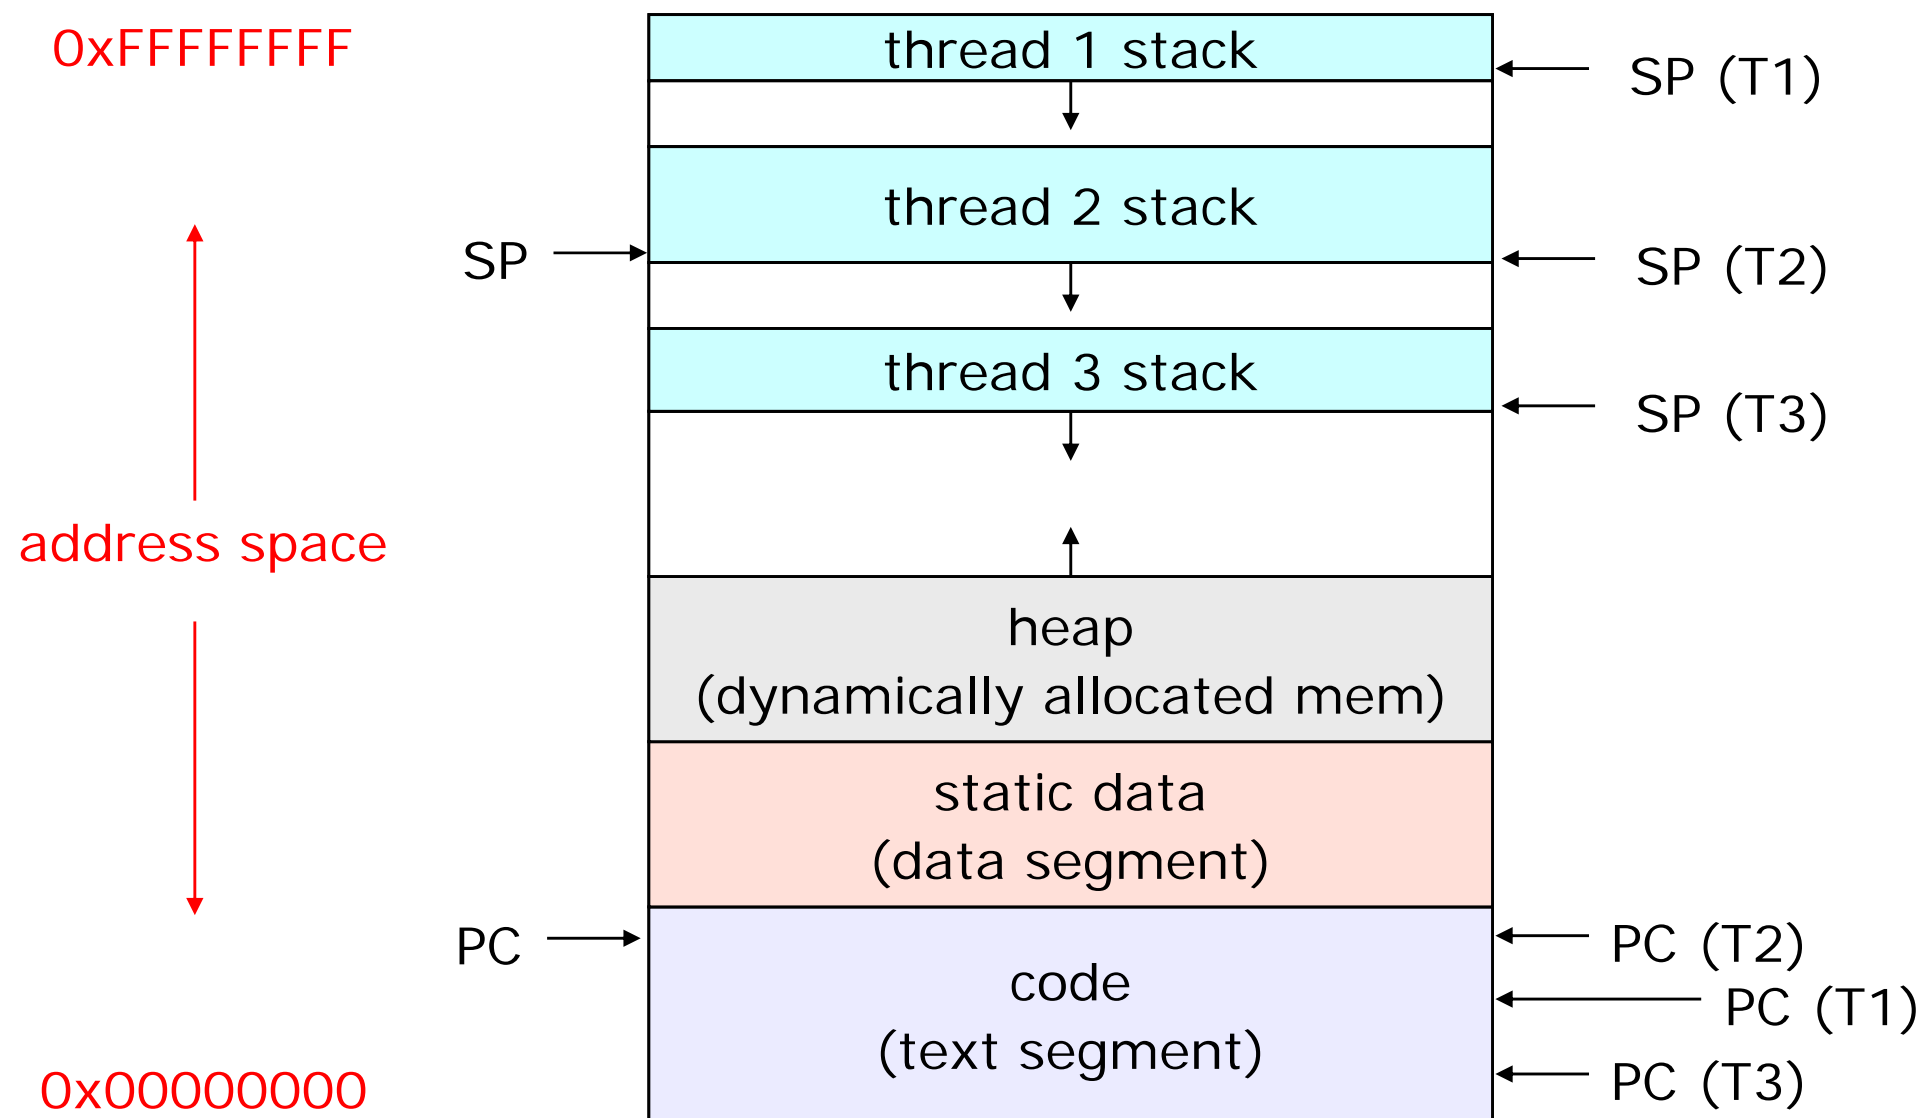
- Separate the concept of a process from its execution state
 - ✓ Process: address space, resources, other general process attributes (e.g., privileges)
 - ✓ Execution state: PC, SP, registers, etc.
 - ✓ This execution state is usually called
 - a thread of control,
 - a thread, or
 - a lightweight process (LWP)



Thread Concept (Cont'd)




Address Space with Threads




Thread Implementation

User-level Threads



Kernel-level Threads



Thread Implementation (Cont'd)

■ User-level threads

- ✓ The user-level threads library implements thread operations
- ✓ They are small and fast
- ✓ User-level threads are invisible to the OS
- ✓ OS may make poor decisions
 - e.g. blocking I/O
- ✓ Thread scheduling
 - Non-preemptive scheduling: `yield()`
 - Preemptive scheduling: timer through signal

■ Kernel-level threads

- ✓ All thread operations are implemented in the kernel
- ✓ The OS schedules all of the threads in a system
- ✓ Kernel threads are cheaper than processes
- ✓ They can still be too expensive



Thread Implementation (Cont'd)

■ Pthreads

- ✓ A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- ✓ API specifies behavior of the thread library, implementation is up to development of the library
- ✓ Common in UNIX operating systems
- ✓ Link with `-lpthread` option

■ Solaris implementation

- ✓ via Light-Weight Process (LWP)
- ✓ Kernel-level implementation (including user-level implementation)

■ Linux implementation

- ✓ Kernel-level implementation, but....
 - a modified process(or task) per thread
- ✓ System call `clone()` for thread creation
- ✓ NGPT (Next Generation POSIX Threading) by IBM



Pthread Libraries for Thread Control

■ Create a thread

- ✓ `#include <pthread.h>`
- ✓ `int pthread_create(pthread_t *tid, pthread_attr_t *attr, void *(start_routine)(void *), void *arg);`
- ✓ return: 0 if OK, nonzero on error

■ Terminate a thread

- ✓ `#include <pthread.h>`
- ✓ `void pthread_exit(void *retval);`

■ Wait for termination of another thread

- ✓ `#include <pthread.h>`
- ✓ `int pthread_join(pthread_t tid, void **tread_return);`
- ✓ return: 0 if OK, nonzero on error



Exercise

■ pthread example

```
$ gcc -o thread thread.c -lpthread (or make thread)
$ ./thread
```

■ Command-line processor: iteration version using one process

```
$ gcc -o cmd_i cmd_i.c (or make cmd_i)
$ ./cmd_i
CMD> doit
Doing doit
Done
CMD> quit
```

■ Command-line processor: a process per command

```
$ gcc -o cmd_p cmd_p.c (or make cmd_p)
$ ./cmd_p
```

■ Command-line processor: a thread per command

```
$ gcc -o cmd_t cmd_t.c -lpthread (or make cmd_t)
$ ./cmd_t
```



Summary

■ System calls in Linux for processes and threads

- ✓ `getpid, getppid, getuid, geteuid, getgid, getegid`
- ✓ `fork`
- ✓ `exit, atexit`
- ✓ `wait, waitpid`
- ✓ `execl, execv, execl, execve, execlp, execvp`
- ✓ `system`
- ✓ `pthread_create, pthread_exit, pthread_join`

