



Inter-Process Communication

경희대학교 컴퓨터공학과

조진성

Inter-Process Communication

Mechanisms for processes to communicate with each other

Linux/Unix IPC

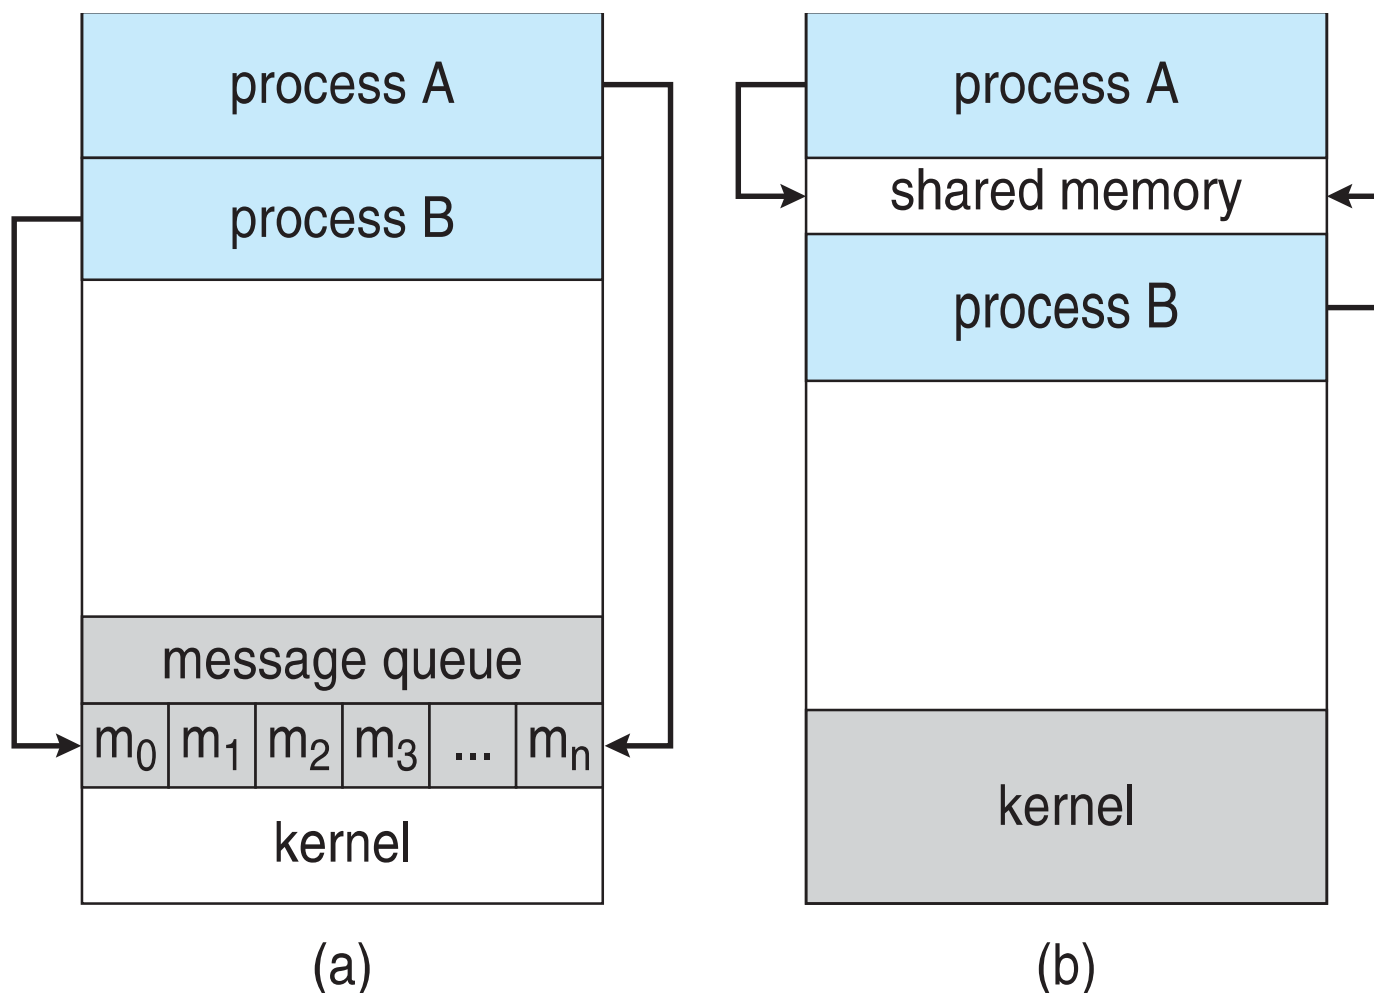
- ✓ pipes
- ✓ FIFOs
- ✓ message queue
- ✓ shared memory
- ✓ sockets



Inter-Process Communication (IPC)

Communication models

- ✓ (a) message passing vs. (b) shared memory



Pipes

The oldest form of UNIX IPC

- ✓ half-duplex
- ✓ between processes that have a common ancestor

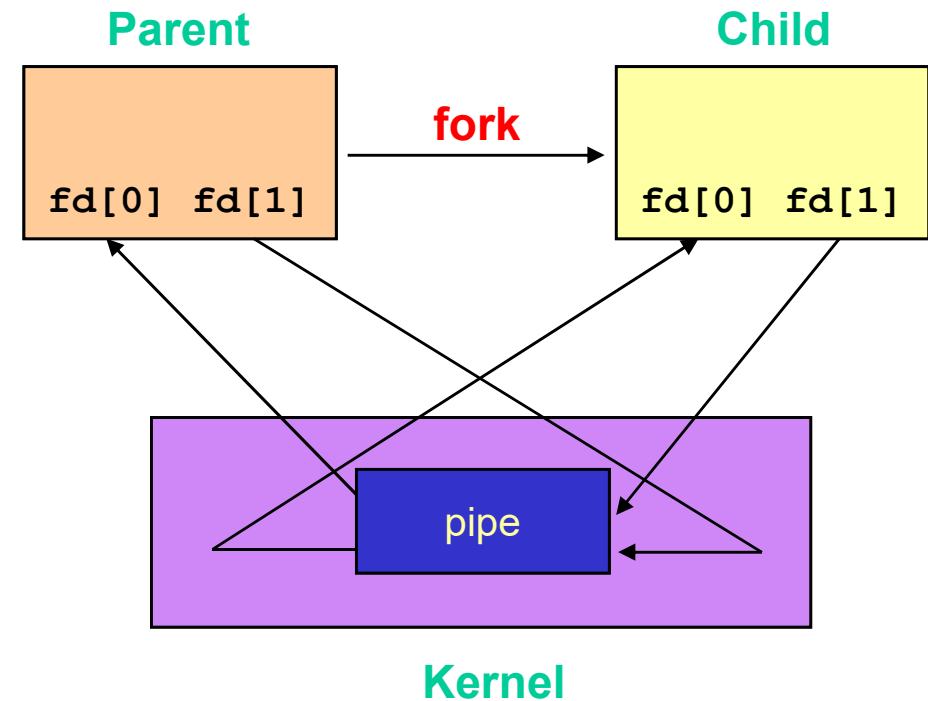
pipe

- ✓ `#include <unistd.h>`
- ✓ `int pipe(int fd[2]);`
- ✓ return: 0 if OK, -1 on error

Half-duplex pipe after a `fork`

IPC through pipes

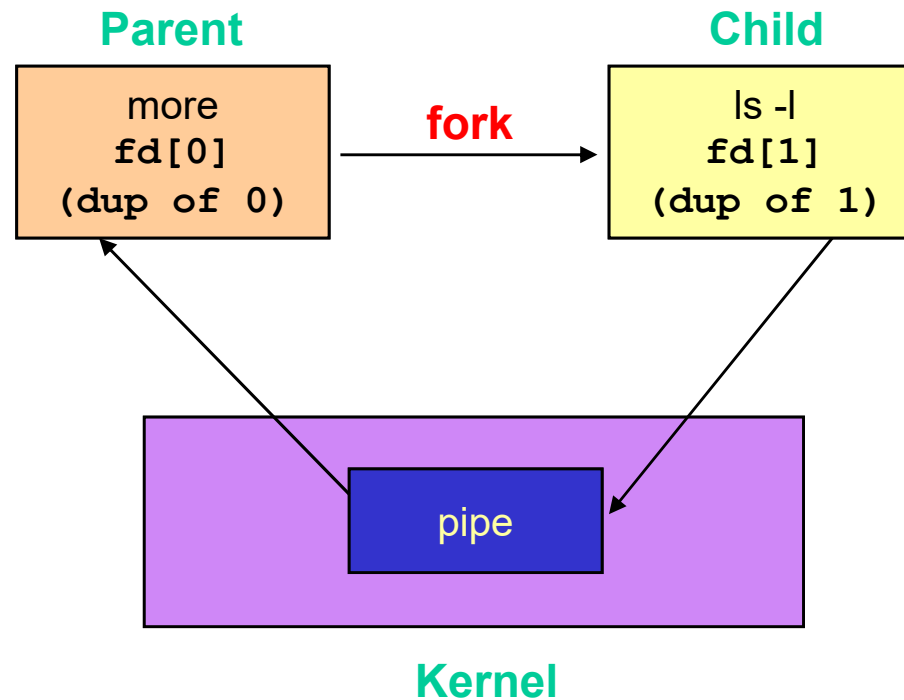
- ✓ Once we have created a pipe using `pipe`,
- ✓ we can use the normal file I/O functions (e.g. `read`, `write`)



Dup

Create a copy of a given file descriptor

- ✓ `#include <unistd.h>`
- ✓ `int dup(int oldfd);`
- ✓ `int dup2(int oldfd, newfd);`
- ✓ return: newfd if OK, -1 on error



Exercise

Send data from parent to child over a pipe

```
$ gcc -o pipe pipe.c (or make pipe)
$ ./pipe
```

Synchronization between parent and child using pipe

```
$ gcc -o sync sync.c synclib.c (or make sync)
$ ./sync
```

Make my own `'ls -l | more'` program using `pipe` & `dup` system call

```
$ gcc -o mymore mymore.c mymore.c (or make mymore)
$ ./mymore
```



FIFOs

Named pipes

- ✓ full-duplex
- ✓ between unrelated processes that don't have a common ancestor

Create a FIFO

- ✓ `#include <sys/types.h>`
- ✓ `#include <sys/stat.h>`
- ✓ `int mkfifo(char *pathname, mode_t mode);`
- ✓ return: 0 if OK, -1 on error
- ✓ Cf) `mkfifo` command

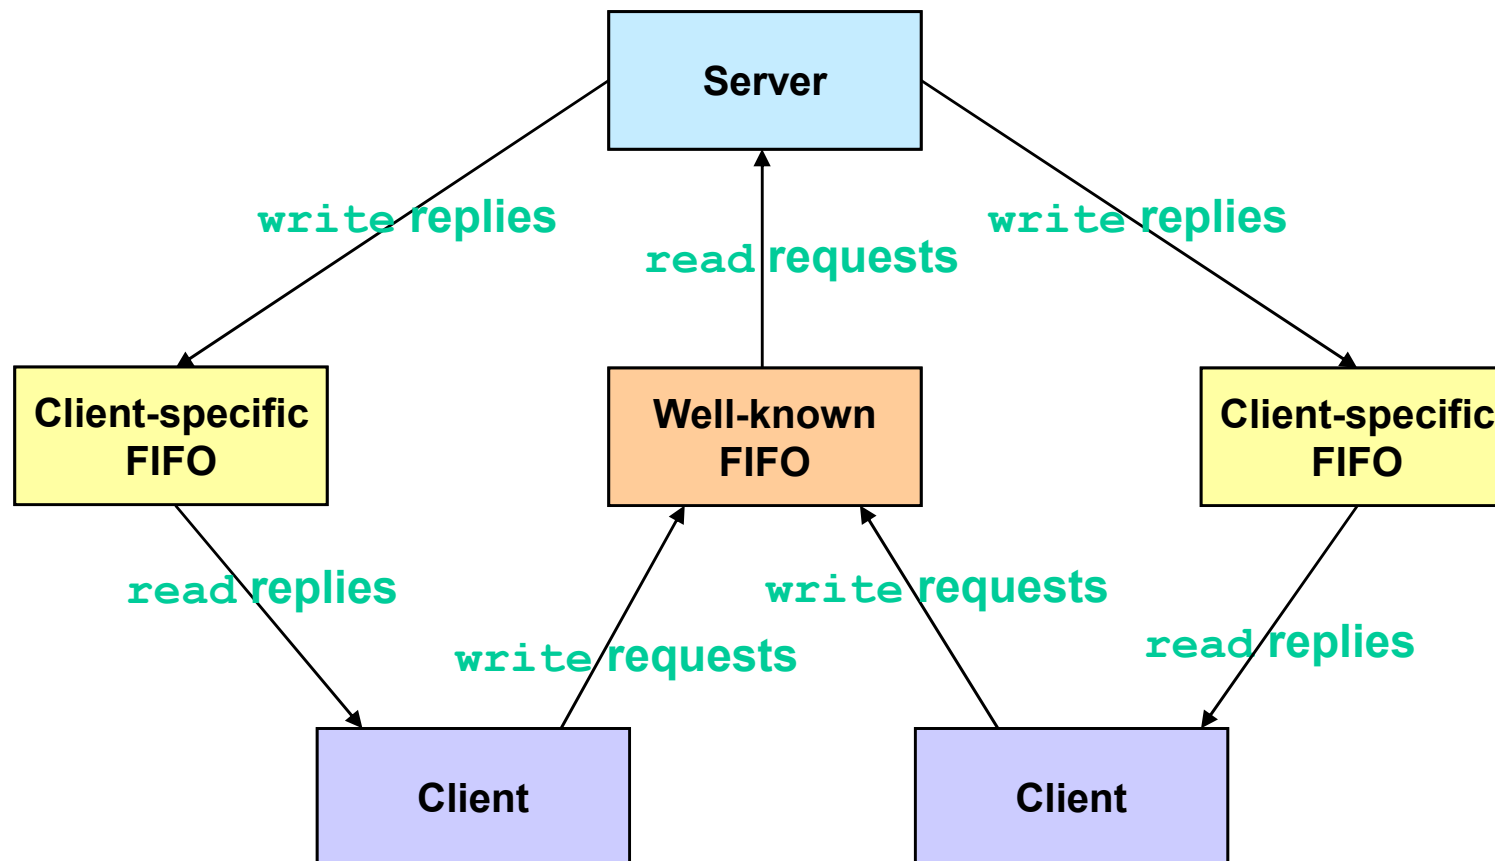
IPC through FIFOs

- ✓ Once we have created a FIFO using `mkfifo`,
- ✓ we can use the normal file I/O functions (e.g. `open`, `read`, `write`, `close`)



FIFOs (Cont'd)

Client-server communication using FIFOs



Exercise

Client-server communication using FIFOs

```
$ gcc -o fifos fifos.c (or make fifos)
```

```
$ gcc -o fifoc fifoc.c (or make fifoc)
```

```
$ ./fifos
```

```
$ ./fifoc
```



Message Queues

A linked list of messages stored within the kernel

A message consists of

- ✓ a long integer that have the positive integer message type
- ✓ message data

```
struct mymsg {  
    long mtype;          /* positive message type */  
    char mtext[512];     /* message data */  
};
```

IPC through message queues

- ✓ **msgget**
 - Open an existing queue or create a new one
- ✓ **msgsnd**
 - Place a message onto the queue
- ✓ **msgrcv**
 - Fetch a message from the queue



System Calls for Message Queues

Obtain a message queue ID

- ✓ `#include <sys/types.h>`
- ✓ `#include <sys/ipc.h>`
- ✓ `#include <sys/msg.h>`
- ✓ `int msgget(key_t key, int flag);`
- ✓ return: message queue ID if OK, -1 on error

Message queue control operations

- ✓ `#include <sys/types.h>`
- ✓ `#include <sys/ipc.h>`
- ✓ `#include <sys/msg.h>`
- ✓ `int msgctl(int msqid, int cmd, struct msqid_ds *buf);`
- ✓ return: 0 if OK, -1 on error



System Calls for Message Queues (Cont'd)

Message queue control operations (Cont'd)

- ✓ The second argument, **cmd**
 - **IPC_STAT** : fetch the **msqid_ds** structure for this queue
 - **IPC_SET** : set the part of **msqid_ds** structure
 - **IPC_RMID** : remove the message queue from the system

- ✓ The third argument, **buf**

```
struct msqid_ds {
    struct ipc_perm  msg_perm;    /* IPC structure: permission and owner */
    struct msg       *msg_first;  /* ptr to first message on queue */
    struct msg       *msg_last;  /* ptr to last message on queue */
    ulong           msg_cbytes;  /* current # of bytes on queue */
    ulong           msg_qnum;    /* # of messages on queue */
    ulong           msg_qbytes;  /* max. # of bytes on queue */
    pid_t           msg_lspid;   /* pid of last msgsnd() */
    pid_t           msg_lrpid;  /* pid of last msgrcv() */
    time_t          msg_stime;   /* last-msgsnd() time */
    time_t          msg_rtime;   /* last-msgrcv() time */
    time_t          msg_ctime;   /* last-change time */
};
```



System Calls for Message Queues (Cont'd)

Send a message

- ✓ `#include <sys/types.h>`
- ✓ `#include <sys/ipc.h>`
- ✓ `#include <sys/msg.h>`
- ✓ `int msgsnd(int msqid, void *ptr, size_t nbytes, int flag);`
- ✓ return: 0 if OK, -1 on error
- ✓ The fourth argument, **flag**
 - `IPC_NOWAIT`

Receive a message

- ✓ `#include <sys/types.h>`
- ✓ `#include <sys/ipc.h>`
- ✓ `#include <sys/msg.h>`
- ✓ `int msgrcv(int msqid, void *ptr, size_t nbytes, long type, int flag);`
- ✓ return: size of data portion of message if OK, -1 on error
- ✓ The fifth argument, **flag**
 - `IPC_NOWAIT`



System Calls for Message Queues (Cont'd)

Receive a message (Cont'd)

- ✓ if **type** == 0,
 - the first message on the queue is returned
- ✓ if **type** > 0,
 - the first message on the queue whose message type equals **type** is returned
- ✓ if **type** < 0,
 - the first message on the queue whose message type is the lowest value less than or equal to the absolute value of **type** is returned
- ✓ if **IPC_NOWAIT** is specified in **flag**,
 - return is made with an error of **EAGAIN**



Exercise

IPC between two processes using message queue

```
$ gcc -o msgq1 msgq1.c (or make msgq1)
```

```
$ gcc -o msgq2 msgq2.c (or make msgq2)
```

```
$ ./msgq1
```

```
$ ./msgq2
```

Note:

- ✓ If a process creates a message queue, its data structure remains in the kernel **even though the process has terminated**
- ✓ You have to remove it through `msgctl()` with `IPC_RMID` parameter
(Or, reboot the system !!!)



Shared Memory

Allows two or more processes to share a given region of memory


- ✓ The fastest form of IPC because data does not need to be copied between them
- ✓ Need to synchronize shared memory access
 - Semaphores are used often

IPC through shared memory

- ✓ **shmget**
 - Obtain a shared memory identifier
 - Open an existing segment or create a new one
- ✓ **shmat**
 - Attach a shared memory segment to the process' address space
- ✓ **shmdt**
 - Detach a shared memory segment
 - **shmdt** does not remove the identifier and its associated data structure from the system

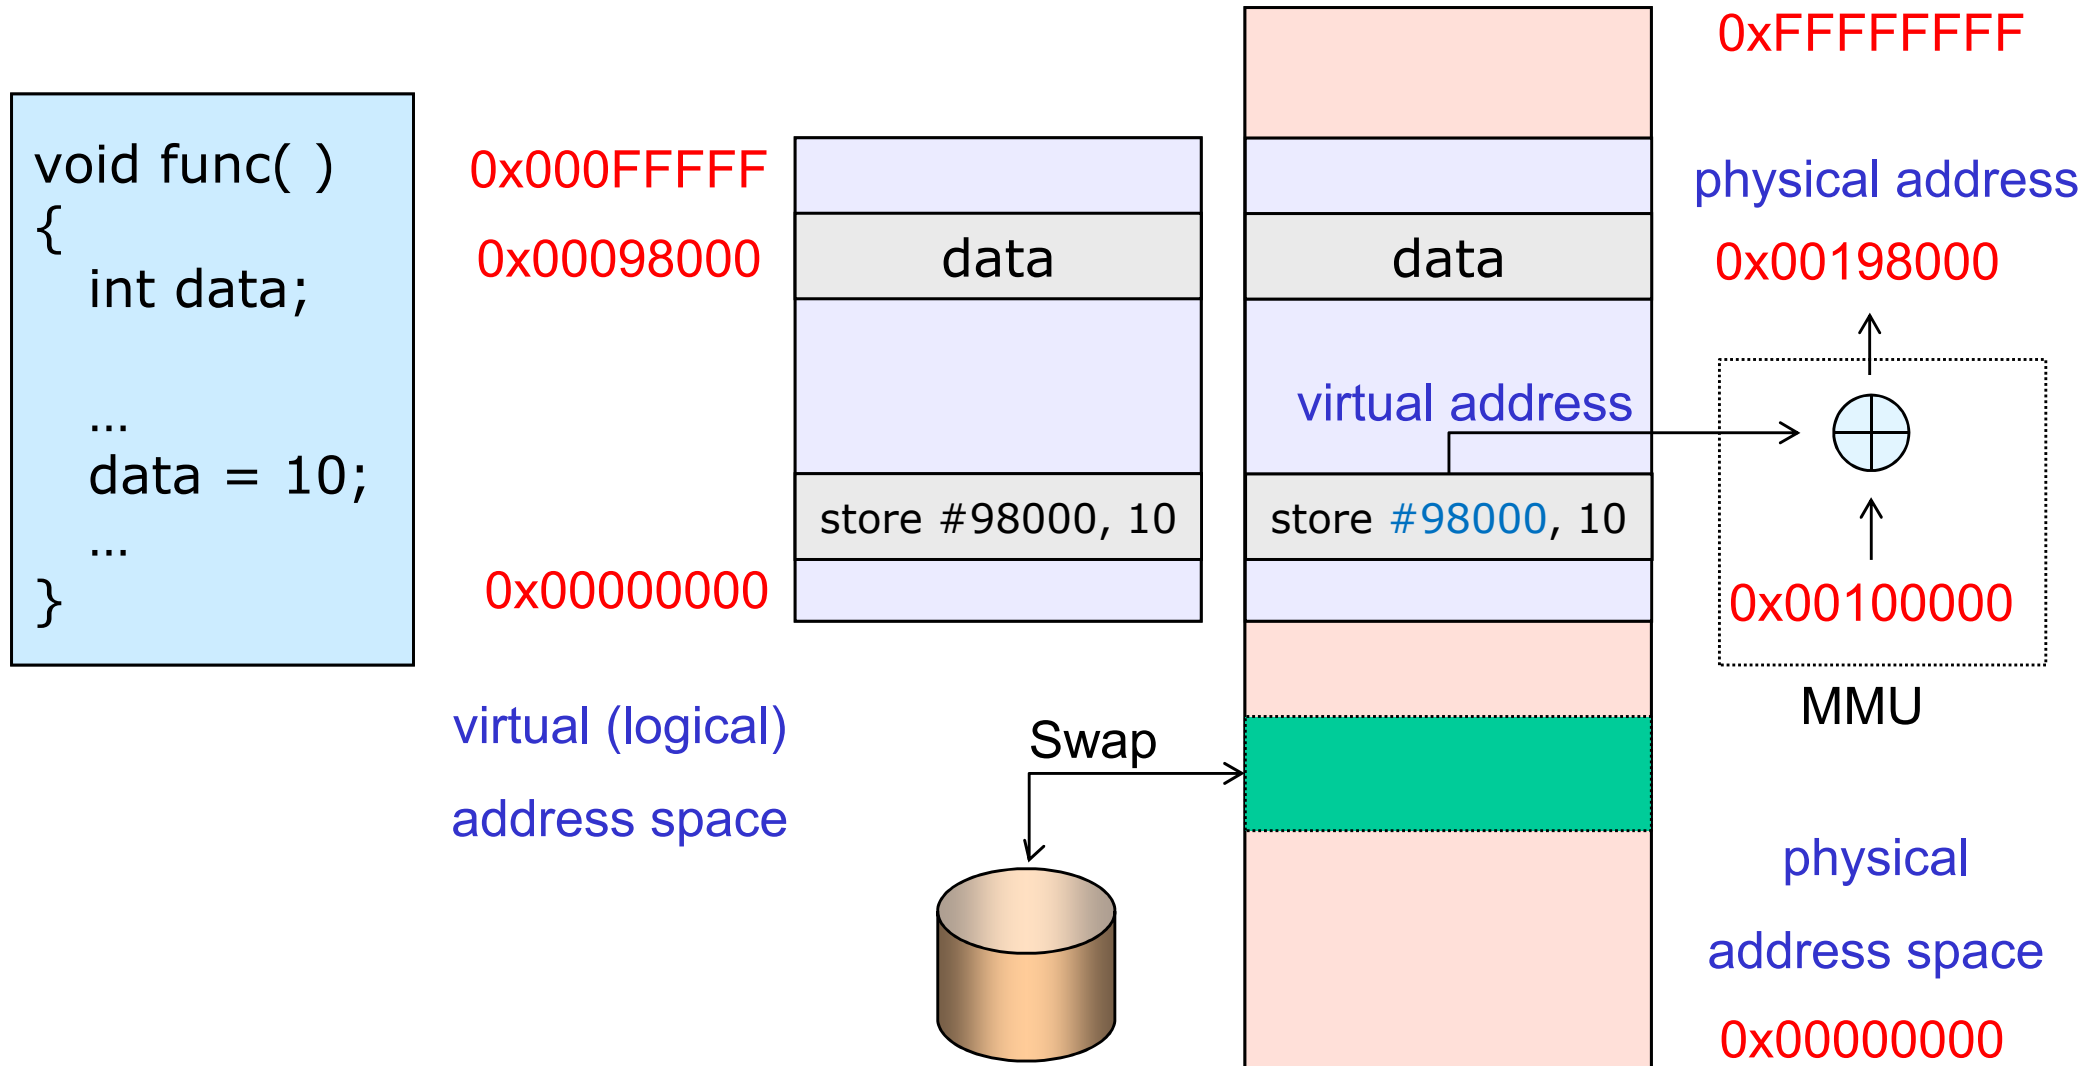


Shared Pages Example



Virtual Memory Management

Address mapping



System Calls for Shared Memory

Obtain a shared memory ID

- ✓ `#include <sys/types.h>`
- ✓ `#include <sys/ipc.h>`
- ✓ `#include <sys/shm.h>`
- ✓ `int shmget(key_t key, int size, int flag);`
- ✓ return: shared memory ID if OK, -1 on error
- ✓ The second argument, **size**
 - if a new segment is being created, we must specify its minimum size, **size**
 - if we are referencing an existing segment, we can specify **size** as 0

Shared memory control operations

- ✓ `#include <sys/types.h>`
- ✓ `#include <sys/ipc.h>`
- ✓ `#include <sys/msg.h>`
- ✓ `int shmctl(int shmid, int cmd, struct shmid_ds *buf);`
- ✓ return: 0 if OK, -1 on error



System Calls for Shared Memory (Cont'd)

Shared memory control operations (Cont'd)

✓ The second argument, **cmd**

- **IPC_STAT** : fetch the **shmid_ds** structure for this shared memory
- **IPC_SET** : set the part of **shmid_ds** structure
- **IPC_RMID** : remove the shared memory segment from the system
- **SHM_LOCK** : lock the shared memory in memory
- **SHM_UNLOCK** : unlock the shared memory segment

✓ The third argument, **buf**

```
struct shmid_ds {
    struct ipc_perm  shm_perm;    /* IPC structure: permission and owner */
    struct anon_map *shm_amp;     /* pointer in kernel */
    int              shm_segsz;   /* size of segment in bytes */
    ushort          shm_lkcnt;   /* # of times segment is being locked */
    pid_t           shm_lpid;    /* pid of last shmop() */
    pid_t           shm_cpid;    /* pid of creator */
    ulong           shm_nattch;  /* # of current attaches */
    ulong           shm_cnattch; /* used only for shminfo */
    time_t          shm_atime;   /* last-attach time */
    time_t          shm_dtime;   /* last-detach time */
    time_t          shm_ctime;   /* last-change time */
};
```



System Calls for Shared Memory (Cont'd)

Attach a shared memory segment

- ✓ `#include <sys/types.h>`
- ✓ `#include <sys/ipc.h>`
- ✓ `#include <sys/shm.h>`
- ✓ `void *shmat(int shmid, void *addr, int flag);`
- ✓ return: pointer to shared memory segment if OK, -1 on error
- ✓ The third argument, **flag**
 - `SHM_RND`, `SHM_RDONLY`
- ✓ if `addr == 0`, (Recommended)
 - the segment is attached at the first available address selected by the kernel
- ✓ if `addr != 0` and `SHM_RND` is not specified,
 - the segment is attached at the address given by `addr`
- ✓ if `addr != 0`, and `SHM_RND` is specified,
 - the segment is attached at the address given by `(addr - (addr modulus SHMLBA))`
- ✓ if `SHM_RDONLY` is specified,
 - the segment is attached read-only



System Calls for Shared Memory (Cont'd)

Detach a shared memory segment

- ✓ `#include <sys/types.h>`
- ✓ `#include <sys/ipc.h>`
- ✓ `#include <sys/shm.h>`
- ✓ `int shmdt(void *addr);`
- ✓ return: 0 if OK, -1 on error



Exercise

Shared memory example (& memory map)

```
$ gcc -o shm shm.c (or make shm)
$ ./shm
```

IPC between two processes using shared memory

```
$ gcc -o sipc1 sipc1.c (or make sipc1)
$ gcc -o sipc2 sipc2.c (or make sipc2)
$ ./sipc1

$ ./sipc2
```

Note:

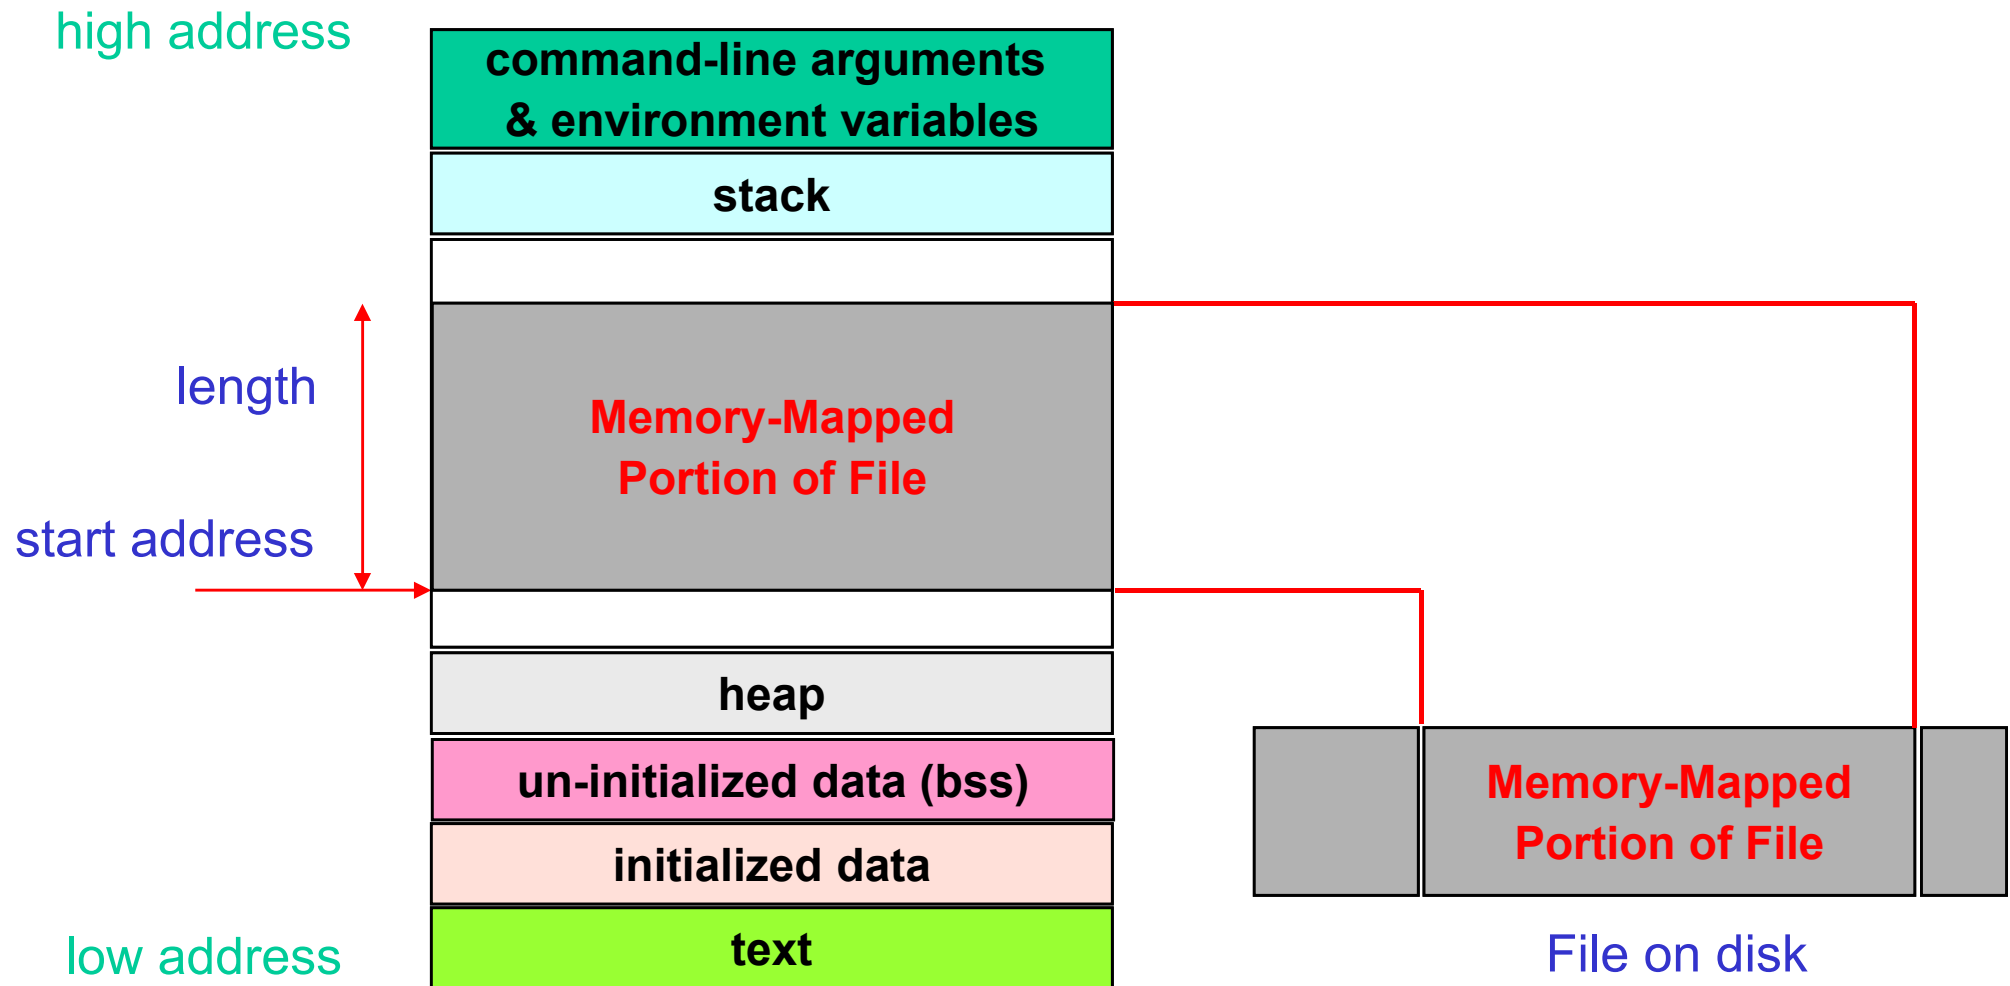
- ✓ If a process creates a shared memory, its data structure remains in the kernel **even though the process has terminated**
- ✓ You have to remove it through `shmctl()` with `IPC_RMID` parameter (Or, reboot the system !!!)



Memory-Mapped File

Map a file on disk into a buffer in memory

- ✓ perform I/O without using `read` or `write`



System Calls for Memory-Mapped File

Map pages of memory

- ✓ `#include <sys/types.h>`
- ✓ `#include <sys/mman.h>`
- ✓ `caddr_t mmap(caddr_t addr, size_t len, int prot, int flag, int fd, off_t off);`
- ✓ return: starting address of mapped region if OK, -1 on error
- ✓ The first argument, **addr**
 - 0 (recommended) : system choose the starting address
 - can be a specific value
- ✓ The third argument, **prot**
 - `PROT_READ` : region can be read
 - `PROT_WRITE` : region can be written
 - `PROT_EXEC` : region can be executed
 - `PROT_NONE` : region cannot be accessed
- ✓ The fourth argument, **flag**
 - `MAP_FIXED` : return value must equal addr
 - `MAP_SHARED` : store operations modify the mapped file
 - `MAP_PRIVATE` : store operations modify a copy of mapped file



System Calls for Memory-Mapped File (Cont'd)

Unmap a memory-mapped region

- ✓ Automatically unmapped when the process terminates, or
- ✓ `#include <sys/types.h>`
- ✓ `#include <sys/mman.h>`
- ✓ `int munmap(caddr_t addr, size_t len);`
- ✓ return: 0 if OK, -1 on error



Exercise

Make my own `cp` program using memory-mapped file

```
$ gcc -o mycp3 mycp3.c (or make mycp3)
```

```
$ ./mycp3 mycp3.c mycp3.bak
```

```
$ ls -l mycp3.c mycp3.bak
```

IPC between parent and child using memory mapped file of `/dev/zero`

```
$ gcc -o mipc mipc.c synclib.c (or make mipc)
```

```
$ ./mipc
```



Summary

Inter-Process Communication (IPC)

- ✓ Mechanisms for processes to communicate with each other

System calls in Linux for IPC

- ✓ Pipe: `pipe`, `read`, `write`, `close`
- ✓ FIFO: `mkfifo`, `open`, `read`, `write`, `close`
- ✓ Message queue: `msgget`, `msgctl`, `msgsnd`, `msgrcv`
- ✓ Shared memory: `shmget`, `shmctl`, `shmat`, `shmdt`
- ✓ Memory-mapped file: `mmap`, `munmap`

