



Synchronization

경희대학교 컴퓨터공학과

조진성

Synchronization

- Threads cooperate in multithreaded programs
 - ✓ To **share** resources, access shared data structures
 - ✓ Also, to **coordinate** their execution

- For correctness, we have to control this cooperation
 - ✓ Must assume threads interleave executions arbitrarily and at different rates
 - Scheduling is not under application writers' control
 - ✓ We control cooperation using **synchronization**
 - Enables us to restrict the interleaving of execution
 - ✓ (Note) This also applies to processes, not just threads
 - And it also applies across machines in a distributed system



An Example

■ Withdraw money from a bank account

- ✓ Suppose you and your girl(boy) friend share a bank account with a balance of 1,000,000won
- ✓ What happens if both go to separate ATM machines, and simultaneously withdraw 100,000won from the account?

```
int withdraw (account, amount)
{
    balance = get_balance (account);
    balance = balance - amount;
    put_balance (account, balance);
    return balance;
}
```



An Example (Cont'd)

■ Interleaved schedules

- ✓ Represent the situation by creating a separate thread for each person to do the withdrawals
- ✓ The execution of the two threads can be interleaved, assuming preemptive scheduling:

**Execution
sequence
as seen by
CPU**

```
balance = get_balance (account);  
balance = balance - amount;
```

**Context
switch**

```
balance = get_balance (account);  
balance = balance - amount;  
put_balance (account, balance);
```

**Context
switch**

```
put_balance (account, balance);
```



Synchronization Problem

■ Problem

- ✓ Two concurrent threads (or processes) access a **shared resource** without any **synchronization**
- ✓ Creates a **race condition**
 - The situation where several processes access and manipulate shared data concurrently
 - The result is non-deterministic and depends on timing
- ✓ **Critical section**
 - Code segment in which the shared data is accessed
- ✓ We need mechanisms for controlling access to critical sections in the face of concurrency
 - So that we can reason about the operation of programs
- ✓ Synchronization is necessary for any shared data structure
 - buffers, queues, lists, etc.




Another Example: Bounded Buffer

■ No synchronization

Producer

```
void producer(data)
{
    while (count==N) ;
    buffer[in] = data;
    in = (in+1) % N;
    count++;
}
```



Consumer

```
void consumer(data)
{
    while (count==0) ;
    data = buffer[out];
    out = (out+1) % N;
    count--;
}
```



Another Example: Bounded Buffer

■ No synchronization: synchronization problem

- ✓ The statement “`count++`” may be implemented in machine language as:

```
register1 = count
register1 = register + 1
count     = register1
```

- ✓ The statement “`count--`” may be implemented as:

```
register2 = count
register2 = register - 1
count     = register2
```

- ✓ Assume `count` is initially 5. One interleaving of statements is:

```
producer: register1 = count           (register1 = 5)
producer: register1 = register1 + 1   (register1 = 6)
consumer: register2 = count           (register2 = 5)
consumer: register2 = register2 - 1   (register2 = 4)
producer: count      = register1      (counter   = 6)
consumer: count      = register2      (counter   = 4)
```

- ✓ The value of `count` may be either 4 or 6, where the correct result should be 5



Exercise

■ Producer & Consumer sharing bounded buffer (multi-threads version)

```
$ gcc -o prodcons_t prodcons_t.c -lpthread (or make prodcons_t)
$ ./prodcons_t
```

■ Producer & Consumer sharing bounded buffer (multi-processes version)

```
$ gcc -o producer producer.c (or make producer)
$ gcc -o consumer consumer.c (or make consumer)
$ ./consumer

$ ./producer
```

■ But,... what about the result?



Synchronization Mechanisms

- Disabling interrupts

- Spinlocks

- ✓ Very primitive, minimal semantics, used to build others
- ✓ Busy waiting

- Semaphores

- ✓ Basic, easy to get the hang of, hard to program with
- ✓ Binary semaphore = mutex (\cong lock)
- ✓ Counting semaphore

- Monitors

- ✓ High-level, requires language support, implicit operations
- ✓ Easy to program with: Java “synchronized”

- Mutex + Condition variables

- ✓ Pthreads



Semaphores

■ Semaphore

- ✓ A counter used to provide access to **a shared data object** for multiple **processes or threads**
- ✓ Two operations
 - wait or P
 - signal or V

■ Synchronization procedure using semaphores

- ✓ Test the semaphore that controls the resource
- ✓ If the value of the semaphore is positive, the process can use the resource
 - The process decrements the semaphore value by 1, indicating that it has used one unit of the resource
- ✓ If the value of the semaphore is 0, the process goes to sleep until the semaphore value is greater than 0
 - When the process wakes up, it returns to above step



Semaphore Implementations

■ System V semaphore

- ✓ Named semaphore → between processes
- ✓ Shared key (number) between processes
- ✓ Serviced by kernel

■ POSIX semaphore

- ✓ Unnamed semaphore → between threads or related processes
- ✓ Shared variable in `sem_t` type between threads or related processes
- ✓ Serviced by libraries or kernel
- ✓ Most implementation doesn't support synchronization between processes yet, including Solaris and Linux



POSIX Semaphores

■ POSIX semaphore libraries

- ✓ `#include <semaphore.h>`
- ✓ `int sem_init(sem_t *sem, int pshared, unsigned int value);`
- ✓ `int sem_wait(sem_t *sem);`
- ✓ `int sem_trywait(sem_t *sem);`
- ✓ `int sem_post(sem_t *sem);`
- ✓ `int sem_getvalue(sem_t *sem, int *sval);`
- ✓ `int sem_destroy(sem_t *sem);`
- ✓ return: 0 if OK, non-zero value on error

■ Note: link option (dependent on semaphore packages)

- ✓ Solaris
 - `-lposix4`
- ✓ Linux
 - `-lpthread`



System Calls for System V Semaphores

■ Obtain a semaphore set ID

- ✓ `#include <sys/types.h>`
- ✓ `#include <sys/ipc.h>`
- ✓ `#include <sys/sem.h>`
- ✓ `int semget(key_t key, int nsems, int flag);`
- ✓ return: semaphore ID if OK, -1 on error

■ Semaphore control operations

- ✓ `#include <sys/types.h>`
- ✓ `#include <sys/ipc.h>`
- ✓ `#include <sys/sem.h>`
- ✓ `int semctl(int semid, int semnum, int cmd, union semun arg);`
- ✓ return: non-negative value depending on `cmd` if OK, -1 on error



System Calls for System V Semaphores (Cont'd)

■ Semaphore control operations (Cont'd)

✓ The third argument, **cmd**

- **IPC_STAT** : fetch the **semid_ds** structure for this semaphore set
- **IPC_SET** : set the part of **semid_ds** structure
- **IPC_RMID** : remove the semaphore set from the system
- **GETVAL** : return the semaphore value
- **SETVAL** : set the semaphore value
- **GETPID** : get pid of the process which do the last access to the semaphore
- **GETNCNT** : return the number of processes which wait for the semaphore to increase
- **GETZCNT** : return the number of processes which wait for the semaphore to be zero
- **GETALL** : fetch all the semaphore values in the set
- **SETALL** : set all the semaphore values in the set

✓ The fourth argument, **arg**

```
union semun {  
    int          val;      /* for SETVAL */  
    struct semid_ds *buf;   /* for IPC_STAT and IPC_SET */  
    ushort       *array;   /* for GETALL and SETALL */  
};
```



System Calls for System V Semaphores (Cont'd)

■ Semaphore operations

- ✓ `#include <sys/types.h>`
- ✓ `#include <sys/ipc.h>`
- ✓ `#include <sys/sem.h>`
- ✓ `int semop(int semid, struct sembuf semop[], size_t nops);`
- ✓ return: 0 if OK, -1 on error
- ✓ The second argument, `semop`

```
struct sembuf {  
    ushort sem_num; /* member # in set (0, 1, ..., nsems-1) */  
    short  sem_op;  /* operation (negative, 0, or positive) */  
    short  sem_flg; /* IPC_NOWAIT, SEM_UNDO */  
};
```



System Calls for System V Semaphores (Cont'd)

■ Semaphore operations (Cont'd)

- ✓ if `sem_op > 0`,
 - the value of `sem_op` is added to the semaphore's value
- ✓ if `sem_op < 0` and the semaphore's value \geq the absolute value of `sem_op`,
 - the value of `sem_op` is added to the semaphore's value
- ✓ if `sem_op < 0` and the semaphore's value $<$ the absolute value of `sem_op`,
 - if `IPC_NOWAIT` is not specified, the calling process is suspended until the semaphore's value \geq the absolute value of `sem_op`
 - if `IPC_NOWAIT` is specified, return is made with an error of `EAGAIN`
- ✓ if `sem_op == 0` and the semaphore's value is not zero,
 - if `IPC_NOWAIT` is not specified, the calling process is suspended until the semaphore's value becomes zero
 - if `IPC_NOWAIT` is specified, return is made with an error of `EAGAIN`

■ Semaphore adjustment on `exit`

- ✓ What happens if a process terminates while it has resources allocated through a semaphore?
- ✓ Use `SEM_UNDO` flags



System Calls for System V Semaphores (Cont'd)

■ POSIX semaphore-like library using System V semaphore

- ✓ `semaphore.h` & `semaphore.c`

- ✓ `int semInit(key_t key);`

- ✓ `int semInitValue(int semid, int value);`

- ✓ `int semWait(int semid);`

- ✓ `int semTryWait(int semid);`

- ✓ `int semPost(int semid);`

- ✓ `int semGetValue(int semid);`

- ✓ `int semDestroy(int semid);`



Exercise

- Implementation of semaphores similar to POSIX semaphores using System V semaphores & shared memory

```
$ gcc -c semlib.c (or make semlib.o)
```

- Note:

- ✓ If a process creates a system V semaphore, its data structure remains in the kernel **even though the process has terminated**
- ✓ You have to remove it through `semctl()` with `IPC_RMID` parameter
(Or, reboot the system !!!)



Exercise (Cont'd)


■ Bounded buffer implementation with semaphores

Producer

```
void produce(data)
{
    wait (empty);
    wait (mutex);
    buffer[in] = data;
    in = (in+1) % N;
    signal (mutex);
    signal (full);
}
```

Semaphore
mutex = 1;
empty = N;
full = 0;

struct item buffer[N];
int in, out;



Consumer

```
void consume(data)
{
    wait (full);
    wait (mutex);
    data = buffer[out];
    out = (out+1) % N;
    signal (mutex);
    signal (empty);
}
```

Exercise

■ Producer & Consumer example using pthreads & POSIX semaphores

```
$ gcc -o prodcons prodcons.c -lpthread  
    (or make prodcons)  
$ ./prodcons
```


■ Producer & Consumer example using `semlib` library

```
$ gcc -o producer_s producer_s.c semlib.c (or make producer_s)  
$ gcc -o consumer_s consumer_s.c semlib.c (or make consumer_s)  
$ ./consumer_s  
  
$ ./producer_s
```



Exercise (Cont'd)

■ Dining philosopher



Exercise (Cont'd)

■ Dining philosopher: A simple solution

```
Semaphore chopstick[N]; // initialized to 1
void philosopher (int i)
{
    while (1) {
        think ();
        wait (chopstick[i]);
        wait (chopstick[(i+1) % N]);
        eat ();
        signal (chopstick[i]);
        signal (chopstick[(i+1) % N]);
    }
}
```

⇒ Problem: causes deadlock



Exercise (Cont'd)

■ Dining philosopher: Deadlock-free version (starvation?)

```
#define N      5
#define L(i)   ((i+N-1)%N)
#define R(i)   ((i+1)%N)
void philosopher (int i) {
    while (1) {
        think ();
        pickup (i);
        eat();
        putdown (i);
    }
}
void test (int i) {
    if (state[i]==HUNGRY &&
        state[L(i)]!=EATING &&
        state[R(i)]!=EATING) {
        state[i] = EATING;
        signal (s[i]);
    }
}
```

```
Semaphore mutex = 1;
Semaphore s[N];
int state[N];

void pickup (int i) {
    wait (mutex);
    state[i] = HUNGRY;
    test (i);
    signal (mutex);
    wait (s[i]);
}
void putdown (int i) {
    wait (mutex);
    state[i] = THINKING;
    test (L(i));
    test (R(i));
    signal (mutex);
}
```



Exercise

■ Dining Philosopher example using pthreads & POSIX semaphores

```
$ gcc -o dining dining.c -lpthread (or make dining)
```

```
$ ./dining
```

✓ Deadlock? Yes? No? Why?

■ Dining Philosopher example using pthreads & POSIX semaphores

✓ Deadlock-free version

```
$ gcc -o dining2 dining2.c -lpthread (or make dining2)
```

```
$ ./dining2
```

✓ Starvation? Yes? No? Why?



Mutexes

■ Mutexes

- ✓ Mutual exclusive locks for threads
- ✓ Serviced by Pthread libraries
- ✓ Similar to binary semaphore

■ Pthread libraries for mutexes

- ✓ `#include <pthread.h>`
- ✓ `int pthread_mutex_init(pthread_mutex_t *mutex,
pthread_mutexattr_t *mattr);`
- ✓ `int pthread_mutex_destroy(pthread_mutex_t *mutex);`
- ✓ `int pthread_mutex_lock(pthread_mutex_t *mutex);`
- ✓ `int pthread_mutex_trylock(pthread_mutex_t *mutex);`
- ✓ `int pthread_mutex_unlock(pthread_mutex_t *mutex);`
- ✓ return: 0 if OK, non-zero value on error



Condition Variables

■ Condition variables

- ✓ A synchronization device that allows threads to suspend and resume execution until some predicate on shared data is satisfied
- ✓ Serviced by Pthread libraries
- ✓ Associated with a mutex, to avoid the race condition

■ Pthread libraries for condition variables

- ✓ `#include <pthread.h>`
- ✓ `int pthread_cond_init(pthread_cond_t *cond,
pthread_condattr_t *cattr);`
- ✓ `int pthread_cond_destroy(pthread_cond_t *cond);`
- ✓ `int pthread_cond_wait(pthread_cond_t *cond,
pthread_mutex_t *mutex);`
- ✓ `int pthread_cond_signal(pthread_cond_t *cond);`
- ✓ `int pthread_cond_broadcast(pthread_cond_t *cond);`
- ✓ return: 0 if OK, non-zero value on error



Exercise

■ Bounded buffer implementation with mutexes and condition variables

```
void producer(data)
{
    while (count==N) ;
    buffer[in] = data;
    in = (in+1) % N;
    count++;
}
```

```
void consumer(data)
{
    while (count==0) ;
    data = buffer[out];
    out = (out+1) % N;
    count--;
}
```

```
pthread_mutex_t mutex;
pthread_cond_t not_full, not_empty;
buffer resources[N];
void producer (resource x) {
    pthread_mutex_lock (&mutex);
    while (array "resources" is full)
        pthread_cond_wait (&not_full, &mutex);
    add "x" to array "resources";
    pthread_cond_signal (&not_empty);
    pthread_mutex_unlock (&mutex);
}
void consumer (resource *x) {
    pthread_mutex_lock (&mutex);
    while (array "resources" is empty)
        pthread_cond_wait (&not_empty, &mutex);
    *x = get resource from array "resources"
    pthread_cond_signal (&not_full);
    pthread_mutex_unlock (&mutex);
}
```



Exercise (Cont'd)

■ Producer & Consumer example using mutexes and condition variables

```
$ gcc -o prodcons_m prodcons_m.c -lpthread  
  (or make prodcons_m)  
$ ./prodcons_m
```

■ Implementation of POSIX semaphores using mutexes and condition variables

```
$ gcc -c semlib2.c (or make semlib2.o)
```

■ Producer & Consumer example using `sem1ib2` library

```
$ gcc -o prodcons_s prodcons_s.c semlib2.c -lpthread  
  (or make prodcons_s)  
$ ./prodcons_s
```



Summary

■ Synchronization

- ✓ Mechanisms for processes and/or threads to control their execution sequences
 - for shared resources in critical section
 - for just execution sequences (e.g. receive before send)

■ System calls in Linux for high-level synchronization mechanisms

- ✓ System V semaphore: `semget`, `semctl`, `semop`
- ✓ POSIX semaphore: `sem_init`, `sem_wait`, `sem_trywait`, `sem_post`,
`sem_getvalue`, `sem_destroy`
- ✓ Mutex: `pthread_mutex_init`, `pthread_mutex_destroy`,
`pthread_mutex_lock`, `pthread_mutex_trylock`,
`pthread_mutex_unlock`
- ✓ Condition variables: `pthread_cond_init`, `pthread_cond_destroy`,
`pthread_cond_wait`, `pthread_cond_signal`,
`pthread_cond_broadcast`

